

**Alex Teboul**  
**CSC 555 Mining Big Data**  
Assignment 4

**Due Friday, March 6<sup>th</sup>**

- 1) Consider a Hadoop job that will result in 87 blocks of output to HDFS.  
Suppose that writing an output block to HDFS takes 1 minute. The HDFS replication factor is set to 3 (for simplicity, we charge reducers for the cost of writing replicated blocks).
- a) How long will it take for the reducer to write the job output on a 5-node Hadoop cluster? (ignoring the cost of Map processing, but counting replication cost in the output writing).
- $87 \text{ blocks} / 5 \text{ nodes} = 17.4 \text{ blocks per node}$
  - 17.4 can round up to ~18 blocks per node
  - $18 \text{ blocks per node} * 3 \text{ replication-factor} = 54 \text{ blocks per node}$
  - $18 \text{ blocks per node} * 1 \text{ min per block per node} = 54 \text{ min}$
  - **~ 54min**
- b) How long will it take for reducer(s) to write the job output to 10 Hadoop worker nodes? (Assume that data is distributed evenly and replication factor is set to 1)
- $87 \text{ blocks} / 10 \text{ nodes} = 8.7 \text{ blocks per node}$
  - 8.7 can round up to ~ 9 blocks per node
  - $9 \text{ blocks per node} * 1 \text{ replication-factor} = 9 \text{ blocks per node}$
  - $9 \text{ blocks per node} * 1 \text{ min per block per node} = 9 \text{ min}$
  - **~ 9 min**
- c) How long will it take for reducer(s) to write the job output to 10 Hadoop worker nodes? (Assume that data is distributed evenly and replication factor is set to 3)
- $87 \text{ blocks} / 10 \text{ nodes} = 8.7 \text{ blocks per node}$
  - 8.7 can round up to ~ 9 blocks per node
  - $9 \text{ blocks per node} * 3 \text{ replication-factor} = 27 \text{ blocks per node}$
  - $27 \text{ blocks per node} * 1 \text{ min per block per node} = 27 \text{ min}$
  - **~ 27 min**
- d) How long will it take for reducer(s) to write the job output to 100 Hadoop worker nodes? (Assume that data is distributed evenly and replication factor is set to 1)
- $87 \text{ blocks} / 100 \text{ nodes} = 0.87 \text{ blocks per node}$
  - 0.87 can round up to ~ 1 blocks per node
  - $1 \text{ blocks per node} * 1 \text{ replication-factor} = 1 \text{ blocks per node}$
  - $1 \text{ blocks per node} * 1 \text{ min per block per node} = 1 \text{ min}$
  - **~ 1 min**


e) How long will it take for reducer(s) to write the job output to 100 Hadoop worker nodes?  
(Assume that data is distributed evenly and replication factor is set to 3)

- $87 \text{ blocks} / 100 \text{ nodes} = 0.87 \text{ blocks per node}$
- $0.87$  can round up to  $\sim 1$  blocks per node
- $1 \text{ blocks per node} * 3 \text{ replication-factor} = 3 \text{ blocks per node}$
- $3 \text{ blocks per node} * 1 \text{ min per block per node} = 3 \text{ min}$
- **$\sim 3 \text{ min}$**

You can ignore the network transfer costs as well as the possibility of node failure.

2)

a) Consider the following graph



Compute the page rank for the nodes in this graph. If you are multiplying matrices manually, you may stop after 5 steps. If you use a tool (e.g., Matlab, website, etc.) for matrix multiplication, you should get your answer to converge.

\*A gets all of B, and a half from Y and Z

\*B gets a half from Y

\*Y gets a half from Z

\*Z gets all of A

**Step 1:** (take transition matrix and multiply by vector  $\rightarrow$  calculate rank)

transition	A	B	Y	Z		V		Rank
A	0	1	1/2	1/2	x	1/4	=	1/2
B	0	0	1/2	0		1/4		1/8
Y	0	0	0	1/2		1/4		1/8
Z	1	0	0	0		1/4		1/4

**Step 2:** (take rank step 1 and use as vector  $\rightarrow$  recalculate rank ... repeat until  $V=Rank$ )

transition	A	B	Y	Z		V		Rank
A	0	1	1/2	1/2	x	1/2	=	5/16
B	0	0	1/2	0		1/8		1/16
Y	0	0	0	1/2		1/8		1/8
Z	1	0	0	0		1/4		1/2

**Step 3:** (... repeat until  $V=Rank$ )

transition	A	B	Y	Z		V		Rank
A	0	1	1/2	1/2	x	5/16	=	3/8
B	0	0	1/2	0		1/16		1/16
Y	0	0	0	1/2		1/8		1/4
Z	1	0	0	0		1/2		5/16

**Step 4:** (... repeat until  $V=Rank$ )

transition	A	B	Y	Z		V		Rank
A	0	1	1/2	1/2	x	3/8	=	11/32
B	0	0	1/2	0		1/16		1/8
Y	0	0	0	1/2		1/4		5/32
Z	1	0	0	0		5/16		3/8

**Step 5:** (... repeat until  $V=Rank$ )

transition	A	B	Y	Z		V		Rank
A	0	1	1/2	1/2	x	11/32	=	25/64
B	0	0	1/2	0		1/8		5/64
Y	0	0	0	1/2		5/32		3/16
Z	1	0	0	0		3/8		11/32

```
[1] import numpy as np
import pandas as pd

[9] M = [(0, 1, 0.5, 0.5),(0,0,0.5,0),(0,0,0,0.5),(1,0,0,0)]
M = np.array(M)

V = [(0.25,),(0.25,),(0.25,),(0.25,)]
V = np.array(V)

print(M, '\n\n', V)

[[[0.  1.  0.5 0.5]
  [0.  0.  0.5 0. ]
  [0.  0.  0.  0.5]
  [1.  0.  0.  0. ]]

  [[0.25]
  [0.25]
  [0.25]
  [0.25]]]

20] V = M.dot(V)
print(V)

[[[0.390625]
  [0.078125]
  [0.1875 ]
  [0.34375 ]]
```

← Step 5 Output

... I tried this in python at it ended up with a rank that seems to be right (order A, B, Y, Z):

```

[[0.36363637]
 [0.09090909]
 [0.18181819]
 [0.36363635]]


```

← After about 20-30ish iterations it starts to converge

+ **Final Page Rank:**


transition	A	B	Y	Z		V		Rank
A	0	1	1/2	1/2	x	4/11	=	4/11
B	0	0	1/2	0		1/11		1/11
Y	0	0	0	1/2		2/11		2/11
Z	1	0	0	0		4/11		4/11

b) Now consider a dead-end node Q and P:



What is the page rank of Q?

\* Following the process from Lecture 7, start with the dead end cut off:



**Step 1:** (take transition matrix and multiply by vector → calculate rank)

transition	A	Y	Z		V		Rank
A	0	1	1/2	x	1/3	=	1/2
Y	0	0	1/2		1/3		1/3
Z	1	0	0		1/3		1/6

....

**Step 30:** it has converged

transition	A	Y	Z		V		Rank
A	0	1	1/2	x	2/5	=	2/5
Y	0	0	1/2		1/5		1/5
Z	1	0	0		2/5		2/5

\* Next, to find the page rank of Q, we plug into a formula the rank according to the ranks of the nodes pointing to Q:

- $\text{PageRank } Q = (1/3) (\text{Rank } Z) + (1/2) (\text{Rank } Y)$
- $\text{PageRank } Q = (1/3) (2/5) + (1/2) (1/5)$
- **PageRank Q = 7/30**       $\sim 0.23333$

Python code for dead-end cut off page calculation:

```
M = [(0, 1, 0.5), (0, 0, 0.5), (1, 0, 0)]
M = np.array(M)

V = [(0.333333333333, ), (0.333333333333, ), (0.333333333333, )]
V = np.array(V)

print(M, '\n\n', V)

[[0.  1.  0.5]
 [0.  0.  0.5]
 [1.  0.  0. ]]

[[0.33333333]
 [0.33333333]
 [0.33333333]]

V = M.dot(V)
print(V)

[[0.39999996]
 [0.20000013]
 [0.40000025]]
```

← After about 20-30ish iterations it starts to converge

c) Exercise 5.1.6 from Mining of Massive Datasets




Figure 5.9: A chain of dead ends

**Exercise 5.1.6:** Suppose we recursively eliminate dead ends from the graph, solve the remaining graph, and estimate the PageRank for the dead-end pages as described in Section 5.1.4. Suppose the graph is a chain of dead ends, headed by a node with a self-loop, as suggested in Fig. 5.9. What would be the PageRank assigned to each of the nodes?

- Going off of the same process as before, if we start with the node that loops to itself, that has a page rank of 1. The second node then gets a rank of  $(1/2) (1) = 1/2$ . The Third node gets a rank of  $(1) (1/2) = 1/2$  because node 2 point only to it. This continues such that all nodes going to the right all keep that rank of  $1/2$ .
- **Page Ranks by Node:**  $1 \rightarrow 1/2 \rightarrow 1/2 \rightarrow \dots \rightarrow 1/2$

3) Given the input data [(1pm, \$6), (2pm, \$15), (3pm, \$15), (4pm, \$20), (5pm, \$10), (6pm, \$20), (7pm, \$20), (8pm, \$24), (9pm, \$23), (10pm, \$30), (11pm, \$30), (12am, \$40)].

a) What will the Hive query “compute average price” return? (yes, this question is as obvious as it seems, asked for comparison with part-b)

- $6+15+15+20+10+20+20+24+23+30+30+40 = 253$
- $253/12 = 21.08$
- **Average price = \$21.08**

b) What will a Storm query “compute average price per each 3 hour window” return? (tumbling, i.e., non-overlapping window of tuples, as many as you can fit). For example, the first window would 1pm-4pm. Second window would be 4pm-7pm.

- **1) 1pm-4pm** (1pm, 2pm, 3pm)
  - $6+15+15 = \$36 \rightarrow \$36/3$
  - **\$12.00**
- **2) 4pm-7pm** (4pm, 5pm, 6pm)
  - $20+10+20 = \$50 \rightarrow \$50/3$
  - **\$16.67**
- **3) 7pm-10pm** (7pm, 8pm, 9pm)
  - $20+24+23 = \$67 \rightarrow \$67/3$
  - **\$22.33**
- **4) 10pm-1am** (10pm, 11pm, 12am)
  - $30+30+40 = \$100 \rightarrow \$100/3$
  - **\$33.33**
- (self-note: \$21.08 avg across windows – all points included)

c) What will a Storm query “compute average price per each 3 hour window” return? (sliding, i.e. overlapping window of tuples, moving the window forward 2 hours each time). First window is 1pm-4pm, second window is 3pm-6pm

- **1) 1pm-4pm** (1pm, 2pm, 3pm)
  - $6+15+15 = \$36 \rightarrow \$36/3$
  - **\$12.00**
- **2) 3pm-6pm** (3pm, 4pm, 5pm)
  - $15+20+10 = \$45 \rightarrow \$45/3$
  - **\$15**
- **3) 5pm-8pm** (5pm, 6pm, 7pm)
  - $10+20+20 = \$50 \rightarrow \$50/3$
  - **\$16.67**
- **4) 7pm-10pm** (7pm, 8pm, 9pm)
  - $20+24+23 = \$67 \rightarrow \$67/3$
  - **\$22.33**
- **5) 9pm-12am** (9pm, 10pm, 11pm)
  - $23+30+30 = \$83 \rightarrow \$83/3$
  - **\$27.67**
- (self-note: \$18.73 avg across windows– missing 12am data point and sliding)

Note, when Storm does not have a full window, you cannot output anything until the window fills with data.

- 4) Run another custom MapReduce job, implementing a solution for the following query:

For Employee(EID, EFirst, ELast, Phone) and Customer(CID, CFirst, CLast, Address), find everyone with the same name using MapReduce:

```
SELECT EFirst, ELast, COUNT(*)
FROM Employee, Customer
WHERE EFirst = CFirst AND ELast = CLast;
GROUP BY EFirst, ELast
```

**\*Needs to include GROUP BY ^**

You can use this input data:

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/employee.txt>  
<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/customer.txt>

```
[ec2-user@ip-172-31-38-169 ~]$ wget http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/employee.txt
--2020-03-07 03:20:37-- http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/employee.txt
Resolving rasinsrv07.cstcis.cti.depaul.edu (rasinsrv07.cstcis.cti.depaul.edu)... 140.192.39.95
Connecting to rasinsrv07.cstcis.cti.depaul.edu (rasinsrv07.cstcis.cti.depaul.edu)|140.192.39.95|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 259133 (253K) [text/plain]
Saving to: 'employee.txt'

employee.txt          100%[=====>] 253.06K  --.-KB/s  in 0.09s

2020-03-07 03:20:38 (2.76 MB/s) - 'employee.txt' saved [259133/259133]

[ec2-user@ip-172-31-38-169 ~]$ wget http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/customer.txt
--2020-03-07 03:21:01-- http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/customer.txt
Resolving rasinsrv07.cstcis.cti.depaul.edu (rasinsrv07.cstcis.cti.depaul.edu)... 140.192.39.95
Connecting to rasinsrv07.cstcis.cti.depaul.edu (rasinsrv07.cstcis.cti.depaul.edu)|140.192.39.95|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 6048428 (5.8M) [text/plain]
Saving to: 'customer.txt'

customer.txt          100%[=====>] 5.77M  9.31MB/s  in 0.6s

2020-03-07 03:21:02 (9.31 MB/s) - 'customer.txt' saved [6048428/6048428]
```

```
ec2-user@ip-172-31-38-169:~
GNU nano 2.5.3 File: customer.txt

800000|Isabella|Cocklin|22 Essex Court, Mount Vernon, NY 10550
800001|Kathryn|Gillmore|535 Cambridge Drive, Lynn, MA 01902
800002|Aja|Pantano|520 Elmwood Avenue, Oshkosh, WI 54901
800003|Isabella|Mcdougale|654 Elmwood Avenue, Highland, IN 46322
800004|Corliss|Favor|960 Winding Way, Millington, TN 38053
800005|Jeanie|Karim|957 Liberty Street, Satellite Beach, FL 32937
800006|Renita|Batey|96 Buttonwood Drive, Port Huron, MI 48060
800007|Genaro|Bagwell|624 Franklin Avenue, Montclair, NJ 07042
```

```
ec2-user@ip-172-31-38-169:~  
GNU nano 2.5.3 File: employee.txt  
EMP0|Francoise|Maynor|72  
EMP1|Isidro|Bosque|80  
EMP2|Brendan|Platt|23  
EMP3|Freedra|Lembke|64  
EMP4|Marcella|Tyrrell|68  
EMP5|Jordan|Mcdougale|78  
EMP6|Sheldon|Read|78
```

Be sure to submit your python code, the command line and the screenshot of successful execution of your code.

### Python Code:

#### Mapper

```
ec2-user@ip-172-31-38-169:~  
GNU nano 2.5.3 File: mapperA4.py  
#!/usr/bin/python  
import sys  
  
# get input  
for line in sys.stdin:  
    #split was on | here  
    line = line.strip()  
    split = line.split('|')  
    #Mapper2 from the employee.txt  
    if split[0].startswith('EMP'):  
        EFirst = split[1]  
        ELast = split[2]  
        FullNameEmployee = EFirst + ' ' + ELast  
        print FullNameEmployee, '\t', 'Employee'  
    #Mapper1 from the customer.txt  
    else:  
        CFirst = split[1]  
        CLast = split[2]  
        FullNameCustomer = CFirst + ' ' + CLast  
        print FullNameCustomer, '\t', 'Customer'  
  
#So I am outputting the full name as will count on FullName Matches.  
█
```



## Reducer

```
ec2-user@ip-172-31-38-169:~  
GNU nano 2.5.3 File: reducerA4.py  
#!/usr/bin/python  
import sys  
  
currentKey = None  
valsEmployee = None  
valsCustomer = None  
  
# input comes from STDIN  
for line in sys.stdin:  
    split = line.strip().split('\t') #[key, value] list  
    key = split[0]  
    value = split[1]  
  
    if currentKey == key: #Same key  
        if value.endswith('Employee'):  
            valsEmployee.append(value)  
        if value.endswith('Customer'):  
            valsCustomer.append(value)  
  
    else:  
        if currentKey:  
            lenEmployee = len(valsEmployee)  
            lenCustomer = len(valsCustomer)  
            if (lenEmployee*lenCustomer > 0):  
                print currentKey, '\t', lenEmployee*lenCustomer  
        currentKey = key  
        if value.endswith('Employee'):  
            valsEmployee = [value]  
            valsCustomer = []  
        elif value.endswith('Customer'):  
            valsEmployee = []  
            valsCustomer = [value]  
lenEmployeeLast = len(valsEmployee)  
lenCustomerLast = len(valsCustomer)  
if (lenEmployeeLast*lenCustomerLast > 0):  
    # join means that there have to be rows on each side  
    print currentKey, '\t', lenEmployeeLast*lenCustomerLast
```

## making directories

```
[ec2-user@ip-172-31-38-169 ~]$ nano mapperA4.py  
[ec2-user@ip-172-31-38-169 ~]$ nano reducerA4.py  
[ec2-user@ip-172-31-38-169 ~]$ hadoop fs -mkdir /data/joinNames  
[ec2-user@ip-172-31-38-169 ~]$ hadoop fs -put customer.txt employee.txt /data/joinNames  
[ec2-user@ip-172-31-38-169 ~]$ hadoop fs -ls /data/joinNames  
Found 2 items  
-rw-r--r--  2 ec2-user supergroup    6048428 2020-03-07 04:22 /data/joinNames/customer.txt  
-rw-r--r--  2 ec2-user supergroup    259133 2020-03-07 04:22 /data/joinNames/employee.txt  
[ec2-user@ip-172-31-38-169 ~]$
```

## copied mapper and reducer over to HADOOP\_HOME to actually run the job.

```
[ec2-user@ip-172-31-38-169 hadoop-2.6.4]$ nano mapperA4.py  
[ec2-user@ip-172-31-38-169 hadoop-2.6.4]$ nano reducerA4.py  
[ec2-user@ip-172-31-38-169 hadoop-2.6.4]$ ls  
bin  etc  include  lib  libexec  LICENSE.txt  logs  mapper03.py  mapperA4.py  NOTICE.txt  README.txt  reducer03.py  reducerA4.py  sbin  share  
[ec2-user@ip-172-31-38-169 hadoop-2.6.4]$
```

## RUN:

```
hadoop jar share/hadoop/tools/lib/hadoop-streaming-2.6.4.jar -input /data/joinNames -mapper  
mapperA4.py -file mapperA4.py -reducer reducerA4.py -file reducerA4.py -output  
/data/outputNames
```

```
can see inside the boxes streamJob logs mapperA4.py mapperA4.py reducerA4.py reducerA4.py reducerA4.py reducerA4.py  
[ec2-user@ip-172-31-38-169 ~]$ hadoop jar share/hadoop/tools/lib/hadoop-streaming-2.6.4.jar -input /data/joinNames -mapper mapperA4.py -file mapperA4.py -r  
-reducerA4.py -file reducerA4.py -output /data/outputNames  
20/03/07 04:41:32 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.  
packageJobJar: [mapperA4.py, reducerA4.py, /tmp/hadoop-unjar2309184568465909972/] [] /tmp/streamjob2270092811295540450.jar tmpDir=null  
20/03/07 04:41:34 INFO client.RMProxy: Connecting to ResourceManager at /172.31.38.169:8032  
20/03/07 04:41:34 INFO client.RMProxy: Connecting to ResourceManager at /172.31.38.169:8032  
20/03/07 04:41:34 INFO mapred.FileInputFormat: Total input paths to process : 2  
20/03/07 04:41:34 INFO mapreduce.JobSubmitter: number of splits:3  
20/03/07 04:41:35 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1583548501086_0013  
20/03/07 04:41:35 INFO impl.YarnClientImpl: Submitted application application_1583548501086_0013  
20/03/07 04:41:35 INFO mapreduce.Job: The url to track the job: http://ip-172-31-38-169.ec2.internal:8088/proxy/application_1583548501086_0013/  
20/03/07 04:41:35 INFO mapreduce.Job: Running job: job_1583548501086_0013  
20/03/07 04:41:41 INFO mapreduce.Job: Job job_1583548501086_0013 running in uber mode : false  
20/03/07 04:41:45 INFO mapreduce.Job: map 0% reduce 0%
```

```
20/03/07 04:41:55 INFO streaming.StreamJob: Output directory: /data/outputNames  
[ec2-user@ip-172-31-38-169 ~]$ hadoop fs -ls /data/outputNames  
Found 2 items  
-rw-r--r--  2 ec2-user supergroup          0 2020-03-07 04:41 /data/outputNames/_SUCCESS  
-rw-r--r--  2 ec2-user supergroup      3924 2020-03-07 04:41 /data/outputNames/part-00000  
[ec2-user@ip-172-31-38-169 ~]$ hadoop fs -cat /data/outputNames/part-00000  
Brendan Anastasio      208  
Brendan Berenbaum      41  
Brendan Bosque         276  
Brendan Cashin         252  
Brendan Lembke         56  
Brendan Mabe           192  
Brendan Maynor         90  
Brendan Mcdougale      102  
Brendan Mullican       196  
Brendan Platt          162  
Brendan Read           378  
Brendan Tyrrell        204  
Brendan Walpole        330  
Francoise Anastasio   343  
Francoise Berenbaum   86  
Francoise Bosque      236  
Francoise Cashin      98  
Francoise Hartley     88  
Francoise Lembke      228  
Francoise Mabe         132  
Francoise Maynor      172  
Francoise Mcdougale    260  
Francoise Mullican    159  
Francoise Platt        46  
Francoise Read         59  
Francoise Tyrrell      51  
Francoise Walpole      176  
Freeda Anastasio       400  
Freeda Berenbaum       240  
Freeda Bosque          268  
Freeda Cashin          51  
Freeda Hartley         196  
Freeda Lembke          372  
Freeda Mabe            132
```

Freeda Maynor	180	
Freeda Mcdougale		275
Freeda Mulligan		345
Freeda Platt	59	
Freeda Read	100	
Freeda Tyrrell		48
Freeda Walpole		220
Hosea Anastasio		350
Hosea Berenbaum		245
Hosea Bosque	147	
Hosea Cashin	177	
Hosea Hartley	290	
Hosea Lembke	328	
Hosea Mabe	220	
Hosea Maynor	192	
Hosea Mcdougale		165
Hosea Mulligan		102
Hosea Platt	141	
Hosea Read	285	
Hosea Tyrrell	84	
Hosea Walpole	165	
Isidro Anastasio		47
Isidro Berenbaum		216
Isidro Bosque	111	
Isidro Cashin	280	
Isidro Hartley		336
Isidro Mabe	132	
Isidro Maynor	258	
Isidro Mcdougale		212
Isidro Mulligan		294

ec2-user@ip-172-31-38-169:~/hadoop-2.6.4

Sheldon Hartley		294
Sheldon Lembke		118
Sheldon Mabe	216	
Sheldon Maynor		216
Sheldon Mcdougale		153
Sheldon Mulligan		192
Sheldon Platt	138	
Sheldon Read	488	
Sheldon Walpole		52
Sherilyn Anastasio		246
Sherilyn Berenbaum		255
Sherilyn Bosque		44
Sherilyn Cashin		82
Sherilyn Hartley		336
Sherilyn Lembke		92
Sherilyn Mabe	141	
Sherilyn Maynor		186
Sherilyn Mcdougale		144
Sherilyn Mulligan		424
Sherilyn Platt		343
Sherilyn Read	56	
Sherilyn Tyrrell		129
Sherilyn Walpole		210
Sid Anastasio	54	
Sid Berenbaum	232	
Sid Bosque	201	
Sid Cashin	153	
Sid Hartley	138	
Sid Lembke	244	
Sid Mabe	159	
Sid Maynor	171	
Sid Mulligan	162	
Sid Platt	245	
Sid Read	47	
Sid Tyrrell	64	
Sid Walpole	364	
Victoria Anastasio		53
Victoria Berenbaum		117
Victoria Bosque		255
Victoria Cashin		245
Victoria Hartley		82
Victoria Lembke		48
Victoria Mabe	180	
Victoria Maynor		200
Victoria Mcdougale		240
Victoria Mulligan		472
Victoria Platt		138
Victoria Read	195	
Victoria Tyrrell		180
Victoria Walpole		64

[ec2-user@ip-172-31-38-169 hadoop-2.6.4]\$

```
Victoria Tyrrell 180
Victoria Walpole 64
[ec2-user@ip-172-31-38-169 ~]$
```

### Sanity Check:

Victoria Walpole = 64 ?

```
EMP150|Sid|Mabe|35
EMP151|Hosea|Walpole|57
EMP152|Victoria|Cashin|50
EMP153|Sheldon|Mulligan|79
EMP154|Victoria|Walpole|53
EMP155|Jordan|Read|34
```

Victoria|Walpole 1/64 ^ v x

^64 instances in the Employee.txt file from

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/employee.txt>

```
820995|Aja|Sano|693 Orchard Street, Algonquin, IL 60102
820996|Giovanni|Smead|992 Oxford Court, Tewksbury, MA 01876
820997|Lavenia|Brock|421 Cypress Court, Schaumburg, IL 60193
820998|Parker|Stephens|655 Park Avenue, Waynesboro, PA 17268
820999|Victoria|Walpole|797 Cedar Street, Muskegon, MI 49441
821000|Ciara|Chupp|128 Summit Avenue, Cranston, RI 02920
821001|Bartul|Bennett|153 Route 2, Sun Prairie, WI 53588
```

Victoria|Walpole 1/1 ^ v x

^1 instance in the Customer.txt from <http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/customer.txt>

➔ I conclude that the MapReduce Job ran effectively given that it worked for Victoria Walpole. Those names were correctly counted. I ran this on the 3-node cluster from Phase1 of the project.

5) In this section you will practice using HBase and setup Mahout and run the curve-clustering example.

- a) Note that HBase runs on top of HDFS, bypassing MapReduce (so only NameNode and DataNode need to be running). You can use your 4-node cluster or the 1-node cluster to run HBase, but be sure to specify which one you used.

### I used the 3-node setup

**cd**

(Download HBase)

**wget http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/hbase-0.90.3.tar.gz**

**gunzip hbase-0.90.3.tar.gz**

**tar xvf hbase-0.90.3.tar**

**cd hbase-0.90.3**

```
hbase-0.90.3/lib/jsr311-api-1.1.1.jar
[ec2-user@ip-172-31-38-169 ~]$ cd hbase-0.90.3
[ec2-user@ip-172-31-38-169 hbase-0.90.3]$
```

(Start HBase service, there is a corresponding stop service and this assumes Hadoop home is set)

**bin/start-hbase.sh**

(Open the HBase shell – at this point jps should show HMaster)

**bin/hbase shell**

```
[ec2-user@ip-172-31-38-169 hbase-0.90.3]$ bin/start-hbase.sh
starting master, logging to /home/ec2-user/hbase-0.90.3/bin/../logs/hbase-0.90.3-169.out
[ec2-user@ip-172-31-38-169 hbase-0.90.3]$ bin/hbase shell
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 0.90.3, r1100350, Sat May  7 13:31:12 PDT 2011

hbase(main):001:0> █
```

(Create an employee table and two column families – private and public. Please watch the quotes, if ' turns into `, the commands will not work)

```
create 'employees', {NAME=> 'private'}, {NAME=> 'public'}
put 'employees', 'ID1', 'private:ssn', '111-222-334'
put 'employees', 'ID2', 'private:ssn', '222-333-445'
put 'employees', 'ID3', 'private:address', '123 State St.'
put 'employees', 'ID1', 'private:address', '243 N. Wabash Av.'
```

**scan 'employees'**

```
hbase(main):001:0> create 'employees', {NAME=> 'private'}, {NAME=> 'public'}
0 row(s) in 0.4250 seconds

hbase(main):002:0> put 'employees', 'ID1', 'private:ssn', '111-222-334'
0 row(s) in 0.0830 seconds

hbase(main):003:0> put 'employees', 'ID2', 'private:ssn', '222-333-445'
0 row(s) in 0.0170 seconds

hbase(main):004:0> put 'employees', 'ID3', 'private:address', '123 State St.'
0 row(s) in 0.0070 seconds

hbase(main):005:0> put 'employees', 'ID1', 'private:address', '243 N. Wabash Av.'
0 row(s) in 0.0080 seconds

hbase(main):006:0> scan 'employees'
ROW
COLUMN+CELL
ID1      column=private:address, timestamp=1583547279391, value=243 N. Wabash Av.
ID1      column=private:ssn, timestamp=1583547239510, value=111-222-334
ID2      column=private:ssn, timestamp=1583547246587, value=222-333-445
ID3      column=private:address, timestamp=1583547260265, value=123 State St.
3 row(s) in 0.0350 seconds
```

Now that we have filled in a couple of values, add 2 new columns to the private family, 1 new column to the public family and create a brand new family with at least 3 columns. For each of these you should introduce at least 2 values -- so a total of  $(2+1+3) * 2 = 12$  values inserted.

➔ add 2 new columns to the private family

```
put 'employees', 'ID1', 'private:school', 'DePaul'
put 'employees', 'ID2', 'private:school', 'Harvard'
```

```
put 'employees', 'ID1', 'private:salary', '66000'  
put 'employees', 'ID2', 'private:salary', '80000'
```

```
hbase(main):007:0> put 'employees', 'ID1', 'private:school', 'DePaul'  
0 row(s) in 0.0070 seconds  
  
hbase(main):008:0> put 'employees', 'ID2', 'private:school', 'Harvard'  
0 row(s) in 0.0060 seconds  
  
hbase(main):009:0> put 'employees', 'ID1', 'private:salary', '66000'  
0 row(s) in 0.0060 seconds  
  
hbase(main):010:0> put 'employees', 'ID2', 'private:salary', '80000'  
0 row(s) in 0.0060 seconds  
  
hbase(main):011:0> scan 'employees'  
ROW COLUMN+CELL  
ID1 column=private:address, timestamp=1583547279391, value=243 N. Wabash Av.  
ID1 column=private:salary, timestamp=1583547425642, value=66000  
ID1 column=private:school, timestamp=1583547406006, value=DePaul  
ID1 column=private:ssn, timestamp=1583547239510, value=111-222-334  
ID2 column=private:salary, timestamp=1583547435927, value=80000  
ID2 column=private:school, timestamp=1583547413256, value=Harvard  
ID2 column=private:ssn, timestamp=1583547246587, value=222-333-445  
ID3 column=private:address, timestamp=1583547260265, value=123 State St.  
3 row(s) in 0.0200 seconds
```

→ 1 new column to the public family

```
put 'employees', 'ID2', 'public:jobtitle', 'data analyst'  
put 'employees', 'ID3', 'public:jobtitle', 'data scientist'
```

```
hbase(main):012:0> put 'employees', 'ID2', 'public:jobtitle', 'data analyst'  
0 row(s) in 0.0070 seconds  
  
hbase(main):013:0> put 'employees', 'ID3', 'public:jobtitle', 'data scientist'  
0 row(s) in 0.0070 seconds  
  
hbase(main):014:0> scan 'employees'  
ROW COLUMN+CELL  
ID1 column=private:address, timestamp=1583547279391, value=243 N. Wabash Av.  
ID1 column=private:salary, timestamp=1583547425642, value=66000  
ID1 column=private:school, timestamp=1583547406006, value=DePaul  
ID1 column=private:ssn, timestamp=1583547239510, value=111-222-334  
ID2 column=private:salary, timestamp=1583547435927, value=80000  
ID2 column=private:school, timestamp=1583547413256, value=Harvard  
ID2 column=private:ssn, timestamp=1583547246587, value=222-333-445  
ID2 column=public:jobtitle, timestamp=1583547479588, value=data analyst  
ID3 column=private:address, timestamp=1583547260265, value=123 State St.  
ID3 column=public:jobtitle, timestamp=1583547486972, value=data scientist  
3 row(s) in 0.0350 seconds
```

→ create a brand new family with at least 3 columns

```
disable 'employees'  
alter 'employees', {NAME=> 'HR'}  
enable 'employees'
```

```
put 'employees', 'ID1', 'HR:complaints', '0'  
put 'employees', 'ID2', 'HR:complaints', '4'
```

```
put 'employees', 'ID1', 'HR:lifeinsurance', 'state farm'  
put 'employees', 'ID2', 'HR:lifeinsurance', 'state farm'
```

```
put 'employees', 'ID1', 'HR:healthinsurance', 'blue cross blue shield'
put 'employees', 'ID2', 'HR:healthinsurance', 'none'
```

```
hbase(main):015:0> disable 'employees'
0 row(s) in 2.0350 seconds

hbase(main):016:0> alter 'employees', {NAME=> 'HR'}
0 row(s) in 0.0340 seconds

hbase(main):017:0> enable 'employees'
0 row(s) in 2.0460 seconds

hbase(main):018:0> put 'employees', 'ID1', 'HR:complaints', '0'
0 row(s) in 0.0070 seconds

hbase(main):019:0> put 'employees', 'ID2', 'HR:complaints', '4'
0 row(s) in 0.0070 seconds

hbase(main):020:0> put 'employees', 'ID1', 'HR:lifeinsurance', 'state farm'
0 row(s) in 0.0060 seconds

hbase(main):021:0> put 'employees', 'ID2', 'HR:lifeinsurance', 'state farm'
0 row(s) in 0.0070 seconds

hbase(main):022:0> put 'employees', 'ID1', 'HR:healthinsurance', 'blue cross blue shield'
0 row(s) in 0.0060 seconds

hbase(main):023:0> put 'employees', 'ID2', 'HR:healthinsurance', 'none'
0 row(s) in 0.0090 seconds
```

Verify that the table has been filled in properly with scan command and submit a screenshot.

```
hbase(main):025:0> scan 'employees'
ROW COLUMN+CELL
ID1 column=HR:complaints, timestamp=1583547711257, value=0
ID1 column=HR:healthinsurance, timestamp=1583547746352, value=blue cross blue shield
ID1 column=HR:lifeinsurance, timestamp=1583547729609, value=state farm
ID1 column=private:address, timestamp=1583547279391, value=243 N. Wabash Av.
ID1 column=private:salary, timestamp=1583547425642, value=66000
ID1 column=private:school, timestamp=1583547406006, value=DePaul
ID1 column=private:ssn, timestamp=1583547239510, value=111-222-334
ID2 column=HR:complaints, timestamp=1583547720728, value=4
ID2 column=HR:healthinsurance, timestamp=1583547754205, value=none
ID2 column=HR:lifeinsurance, timestamp=1583547739773, value=state farm
ID2 column=private:salary, timestamp=1583547435927, value=80000
ID2 column=private:school, timestamp=1583547413256, value=Harvard
ID2 column=private:ssn, timestamp=1583547246587, value=222-333-445
ID2 column=public:jobtitle, timestamp=1583547479588, value=data analyst
ID3 column=private:address, timestamp=1583547260265, value=123 State St.
ID3 column=public:jobtitle, timestamp=1583547486972, value=data scientist
3 row(s) in 0.0410 seconds
```

➔ worked

b) Download and setup Mahout:

**cd**

(download mahout zip package)

**wget <http://rasinrv07.csteis.cti.depaul.edu/CSC555/apache-mahout-distribution-0.11.2.zip>**  
(Unzip the file)



## unzip apache-mahout-distribution-0.11.2.zip

set the environment variables (as always, you can put these commands in ~/.bashrc to automatically set these variables every time you open a new connection, source ~/.bashrc to refresh)

**export MAHOUT\_HOME=/home/ec2-user/apache-mahout-distribution-0.11.2**

**export PATH=/home/ec2-user/apache-mahout-distribution-0.11.2/bin:\$PATH**

be absolutely sure you set Hadoop home variable (if you haven't):

Download and prepare synthetic data – it represents a list of 2D curves, represented as a 50-point vector.

Download the synthetic data example:

**wget [http://rasinsrv07.csteis.eti.depaul.edu/CSC555/synthetic\\_control.data](http://rasinsrv07.csteis.eti.depaul.edu/CSC555/synthetic_control.data)**

(make a testdata directory in HDFS, the example KMeans algorithm assumes the data lives there by default.

\*

start-dfs.sh

start-yarn.sh

mr-jobhistory-daemon.sh start historyserver

jps

```
[ec2-user@ip-172-31-38-169 ~]$ jps
3318 Jps
2809 HMaster
[ec2-user@ip-172-31-38-169 ~]$ start-dfs.sh
Starting namenodes on [ip-172-31-38-169.ec2.internal]
ip-172-31-38-169.ec2.internal: starting namenode, logging to /home/ec2-user/hadoop-2.6.4/logs/hadoop-ec2-user-namenode-ip-172-31-38-169.out
172.31.38.169: starting datanode, logging to /home/ec2-user/hadoop-2.6.4/logs/hadoop-ec2-user-datanode-ip-172-31-38-169.out
172.31.34.17: starting datanode, logging to /home/ec2-user/hadoop-2.6.4/logs/hadoop-ec2-user-datanode-ip-172-31-34-17.out
172.31.45.120: starting datanode, logging to /home/ec2-user/hadoop-2.6.4/logs/hadoop-ec2-user-datanode-ip-172-31-45-120.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /home/ec2-user/hadoop-2.6.4/logs/hadoop-ec2-user-secondarynamenode-ip-172-31-38-169.out
[ec2-user@ip-172-31-38-169 ~]$ start-yarn.sh
starting yarn daemons
starting resourcemanager, logging to /home/ec2-user/hadoop-2.6.4/logs/yarn-ec2-user-resourcemanager-ip-172-31-38-169.out
172.31.45.120: starting nodemanager, logging to /home/ec2-user/hadoop-2.6.4/logs/yarn-ec2-user-nodemanager-ip-172-31-45-120.out
172.31.34.17: starting nodemanager, logging to /home/ec2-user/hadoop-2.6.4/logs/yarn-ec2-user-nodemanager-ip-172-31-34-17.out
172.31.38.169: starting nodemanager, logging to /home/ec2-user/hadoop-2.6.4/logs/yarn-ec2-user-nodemanager-ip-172-31-38-169.out
[ec2-user@ip-172-31-38-169 ~]$ mr-jobhistory-daemon.sh start historyserver
starting historyserver, logging to /home/ec2-user/hadoop-2.6.4/logs/mapred-ec2-user-historyserver-ip-172-31-38-169.out
[ec2-user@ip-172-31-38-169 ~]$ jps
4348 Jps
2809 HMaster
3867 ResourceManager
3432 NameNode
3970 NodeManager
4311 JobHistoryServer
3721 SecondaryNameNode
3560 DataNode
[ec2-user@ip-172-31-38-169 ~]$
```

## hadoop fs -mkdir -p testdata

(copy the synthetic data over to the testdata directory on HDFS side. You can inspect the contents of the file by running **nano synthetic\_control.data** – as you can see this is a list of 600 vectors, with individual values separated by a space)

```
ec2-user@ip-172-31-38-169:~
GNU nano 2.5.3 File: synthetic_control.data
8.7812 34.4632 31.3381 31.2834 28.9207 33.7596 25.3969 27.7849 35.2479 27.1159 32.8717 29.2171 36.0253 32.337 34.5249
4.8923 25.741 27.5532 32.8217 27.8789 31.5926 31.4861 35.5469 27.9516 31.6595 27.5415 31.1887 27.4867 31.391 27.811
1.3987 30.6316 26.3983 24.2905 27.8613 28.5491 24.9717 32.4358 25.2239 27.3068 31.8387 27.2587 28.2572 26.5819 24.0455
5.774 30.5262 35.4209 25.6033 27.97 25.2702 28.132 29.4268 31.4549 27.32 28.9564 28.9916 29.9578 30.2773 30.4447
7.1798 29.2498 33.6928 25.6264 24.6555 28.9446 35.798 34.9446 24.5596 34.2366 27.9634 25.3216 35.4154 34.862 25.1472
5.5067 29.7929 28.0765 34.4812 33.8 27.6671 30.6122 25.6393 30.1171 26.5188 30.1524 27.8514 29.5582 32.3601 29.2064
```

**hadoop fs -put synthetic\_control.data testdata/**



```
[ec2-user@ip-172-31-38-169 ~]$ hadoop fs -mkdir -p testdata
[ec2-user@ip-172-31-38-169 ~]$ hadoop fs -put synthetic_control.data testdata/
```

Please be sure to report the runtime of any command that includes “time”

**time mahout org.apache.mahout.clustering.syntheticcontrol.kmeans.Job**

```
20/03/07 02:44:47 INFO ClusterDumper: Wrote 6 clusters
20/03/07 02:44:47 INFO MahoutDriver: Program took 221756 ms (Minutes: 3.6959333333333335)

real    3m47.223s
user    0m13.620s
sys     0m3.247s
[ec2-user@ip-172-31-38-169 ~]$
```

```
[ec2-user@ip-172-31-38-169 ~]$ hadoop fs -ls output
Found 15 items
-rw-r--r--    2 ec2-user  supergroup          194 2020-03-07 02:44 output/_policy
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:44 output/clusteredPoints
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:41 output/clusters-0
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:41 output/clusters-1
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:44 output/clusters-10-final
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:42 output/clusters-2
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:42 output/clusters-3
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:42 output/clusters-4
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:43 output/clusters-5
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:43 output/clusters-6
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:43 output/clusters-7
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:43 output/clusters-8
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:44 output/clusters-9
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:41 output/data
drwxr-xr-x    - ec2-user  supergroup           0 2020-03-07 02:41 output/random-seeds
```

^it worked

(clusterdump is a built-in Mahout command that will produce the result of KMeans. Output file is written to clusters-10-final because that is where the output is written after 10 iterations. The center points are placed in a separate file, called clusteredPoints)

**mahout clusterdump --input output/clusters-10-final --pointsDir output/clusteredPoints --output clusteranalyze.txt**

```
[ec2-user@ip-172-31-38-169 ~]$ mahout clusterdump --input output/clusters-10-final --pointsDir output/clusteredPoints --output clusteranalyze.txt
Running on hadoop, using /home/ec2-user/hadoop-2.6.4/bin/hadoop and HADOOP_CONF_DIR=
MAHOUT-JOB: /home/ec2-user/apache-mahout-distribution-0.11.2/mahout-examples-0.11.2-job.jar
20/03/07 02:54:06 INFO AbstractJob: Command line arguments: [--dictionaryType=[text], --distanceMeasure=[org.apache.mahout.common.distance.SquaredEuclideanDistanceMeasure], --endPhase=[2147483647], --input=[output/clusters-10-final], --output=[clusteranalyze.txt], --outputFormat=[TEXT], --pointsDir=[output/clusteredPoints], --startPhase=[0], --tempDir=[temp]]
20/03/07 02:54:08 INFO ClusterDumper: Wrote 6 clusters
20/03/07 02:54:08 INFO MahoutDriver: Program took 1710 ms (Minutes: 0.0285)
[ec2-user@ip-172-31-38-169 ~]$
```

The file clusteranalyze.txt contains the results of the Kmeans run after 10 iterations.

Submit the screenshot of the first page from clusteranalyze.txt (e.g., from more clusteranalyze.txt)

**more clusteranalyze.txt**

ec2-user@ip-172-31-38-169:~

```
bioproject.xml          dwddate.tbl            lineorder.tbl          pig-0.15.0.tar
[ec2-user@ip-172-31-38-169 ~]$ more clusteranalyze.txt
{"r": [3.659, 3.666, 2.911, 3.336, 3.885, 3.409, 3.564, 3.38, 3.855, 3.895, 3.53, 3.637, 4.173, 4.063, 4.345, 4.074, 4.733, 4.385, 5.028, 4.924, 4.848,
5.334, 5.298, 5.706, 5.291, 5.469, 5.038, 4.717, 5.084, 4.415, 3.761, 4.362, 4.206, 3.732, 4.083, 3.938, 4.098, 4.417, 3.967, 4.141, 3.814, 4.133, 3.
811, 4.001, 3.885, 4.355, 4.104, 4.182, 4.403, 4.234, 4.574, 4.158, 4.448, 4.876, 4.456, 4.409, 4.974, 4.689, 5.208, 4.925], "c": [30.125, 29.706, 30.
494, 31.896, 31.445, 32.689, 31.704, 32.251, 32.654, 32.738, 33.356, 33.61, 33.88, 34.146, 34.028, 34.665, 35.013, 35.707, 35.907, 35.846, 37.208, 3
9.062, 38.992, 39.687, 40.511, 41.14, 42.314, 41.279, 42.084, 43.498, 43.04, 44.39, 43.608, 43.837, 44.926, 45.259, 45.269, 45.376, 45.56, 46.084, 4
6.632, 47.08, 47.452, 47.246, 47.82, 48.425, 48.837, 48.397, 48.299, 49.408, 49.318, 49.718, 50.716, 50.739, 50.567, 51.405, 51.497, 51.501, 51.79,
51.925], "n": 76, "identifier": "VL-454"}
Weight : [props - optional]: Point:
1.0 : [distance=40.121510962400315]: [33.786, 29.428, 27.377, 37.342, 26.013, 36.203, 31.024, 37.785, 35.754, 36.953, 32.401, 37.258,
37.536, 40.414, 33.863, 33.254, 43.233, 40.022, 34.117, 38.453, 42.565, 37.477, 37.266, 43.044, 39.428, 43.563, 45.807, 49.265, 40.014, 43.101, 40
.967, 47.69, 47.25, 49.214, 43.991, 51.439, 49.912, 53.279, 45.193, 45.813, 48.885, 55.934, 50.823, 53.761, 48.767, 57.682, 58.6, 54.354, 57.292, 50
.088, 55.553, 50.929, 56.859, 54.562, 53.578, 62.184, 62.853, 57.198, 63.828, 56.127]
1.0 : [distance=32.790381834986235]: [24.944, 31.231, 27.187, 29.492, 32.562, 27.971, 32.59, 32.216, 34.967, 40.144, 37.822, 37.343,
36.915, 32.309, 39.231, 36.764, 38.671, 33.74, 42.31, 34.92, 39.418, 39.438, 45.279, 39.32, 43.148, 36.097, 36.905, 44.596, 47.829, 42.111, 39.428,
46.569, 40.807, 44.658, 45.115, 44.59, 41.423, 47.835, 53.857, 44.565, 45.324, 54.725, 44.234, 44.977, 55.838, 52.971, 54.403, 54.437, 56.912, 55.3
66, 49.803, 50.875, 52.817, 60.459, 54.858, 51.21, 53.154, 56.914, 57.193, 57.238]
1.0 : [distance=35.44696249354358]: [24.309, 25.021, 28.828, 36.801, 35.288, 31.411, 36.021, 31.506, 28.032, 32.399, 40.066, 30.909,
40.848, 39.512, 42.186, 38.253, 37.241, 31.94, 32.407, 41.97, 38.037, 43.011, 38.482, 38.847, 35.326, 46.824, 40.332, 39.543, 43.595, 47.021, 48.15
9, 40.292, 47.752, 41.687, 46.825, 46.896, 46.978, 49.747, 44.497, 50.599, 50.357, 53.346, 48.471, 54.013, 54.785, 53.627, 49.751, 46.669, 53.625, 5
8.284, 55.884, 52.417, 58.463, 53.906, 61.116, 52.002, 50.094, 58.254, 53.513, 59.472]
1.0 : [distance=29.864710384679814]: [24.37, 27.876, 30.794, 36.752, 30.161, 34.208, 30.799, 37.802, 33.658, 37.964, 34.201, 31.489,
33.074, 29.063, 39.339, 40.322, 34.684, 38.019, 41.299, 41.739, 31.535, 38.859, 35.895, 35.926, 42.162, 33.023, 34.18, 33.841, 39.712, 40.885, 41.2
69, 44.043, 41.098, 46.379, 43.458, 46.668, 42.184, 42.53, 41.612, 45.165, 44.564, 42.871, 42.466, 39.759, 40.43, 50.122, 45.812, 49.335, 44.666, 51
.675, 47.506, 49.791, 52.686, 45.285, 48.131, 51.817, 49.261, 44.194, 51.162, 53.758]
1.0 : [distance=30.522592692742258]: [32.131, 31.711, 35.793, 27.419, 27.617, 33.054, 36.064, 36.964, 29.058, 30.386, 38.765, 32.313,
32.508, 38.514, 32.682, 37.712, 35.532, 40.359, 39.945, 44.03, 44.16, 39.778, 37.261, 42.001, 40.865, 38.836, 43.657, 37.247, 43.264, 43.675, 38.1
95, 45.961, 40.361, 40.706, 50.959, 52.071, 44.35, 47.271, 42.643, 45.087, 44.759, 46.139, 48.421, 47.248, 54.849, 52.391, 57.017, 56.663, 49.071, 4
9.841, 52.534, 55.698, 53.699, 57.587, 50.934, 59.454, 57.07, 51.82, 53.359, 52.761]
1.0 : [distance=33.628914760843266]: [35.3, 32.696, 26.872, 36.786, 29.216, 33.532, 38.174, 34.597, 27.976, 33.232, 31.12, 37.324, 39
.475, 40.481, 32.965, 37.926, 35.162, 43.285, 41.715, 41.986, 41.577, 42.711, 34.716, 37.45, 40.135, 36.166, 44.384, 39.674, 38.027, 40.841, 41.242
, 40.714, 41.372, 49.402, 44.293, 46.714, 40.813, 44.64, 42.461, 44.351, 53.529, 49.409, 50.862, 44.102, 55.099, 49.759, 54.1, 53.312, 49.772, 56.38
5, 52.057, 54.642, 52.383, 58.083, 56.634, 59.473, 50.001, 60.121, 50.345, 59.133]
1.0 : [distance=25.678636934660688]: [29.74, 25.589, 27.592, 26.803, 26.792, 37.297, 31.595, 35.279, 35.941, 31.392, 34.278, 36.094,
35.631, 40.259, 41.36, 34.375, 33.874, 35.874, 34.566, 37.009, 39.506, 41.6, 40.852, 38.998, 38.457, 36.436, 42.124, 43.112, 38.568, 38.653, 44.7, 4
5.539, 47.969, 39.801, 44.165, 46.559, 38.537, 43.642, 48.437, 43.012, 48.225, 49.461, 51.13, 49.544, 52.623, 41.796, 45.836, 45.249, 48.35, 54.633
, 45.594, 49.112, 50.814, 47.246, 53.784, 52.405, 51.871, 56.813, 46.988, 47.553]
1.0 : [distance=28.781611690879696]: [33.732, 28.655, 25.656, 31.965, 28.901, 31.412, 33.624, 32.795, 37.216, 36.401, 35.106, 37.772,
38.784, 32.891, 38.877, 36.701, 33.854, 32.929, 36.603, 37.34, 32.838, 42.607, 36.403, 38.414, 36.486, 38.804, 43.887, 39.881, 36.683, 40.521, 36.
163, 38.18, 38.217, 45.054, 40.497, 44.442, 47.772, 38.464, 39.243, 49.324, 42.545, 43.105, 46.369, 44.655, 45.206, 50.669, 49.07, 49.605, 40.852, 4
4.198, 47.569, 47.022, 43.183, 53.764, 44.984, 54.396, 45.536, 52.173, 53.873, 54.602]
1.0 : [distance=35.51648243457966]: [35.432, 30.61, 26.088, 25.674, 29.148, 35.054, 31.818, 36.47, 30.972, 30.353, 34.86, 39.275, 39.
--More-- (1%)
```

```
.72, 50.542, 45.538, 45.869, 52.509, 46.161, 46.877, 56.049, 52.835, 54.904, 55.355, 5
1.0 : [distance=32.28028917196654]: [35.939, 26.784, 26.775, 35.103, 36
.441, 36.404, 33.796, 34.848, 42.91, 34.183, 44.889, 36.443, 44.744, 39.13, 43.207, 40
1.089, 47.043, 44.158, 53.626, 45.012, 50.85, 57.246, 48.328, 48.739, 50.743, 55.603,
1.0 : [distance=30.729118469539213]: [29.686, 25.753, 28.764, 35.712, 2
[ec2-user@ip-172-31-38-169 ~]$
```

➔ worked

Submit a single document containing your written answers. Be sure that this document contains your name and “CSC 555 Assignment 4” at the top.